

Image Segmentation

Introduction

Image segmentation is a fundamental task in computer vision that involves partitioning an image into meaningful regions, often corresponding to objects or specific structures. In this Task, we work with the **CAMVid dataset**, a labeled dataset for semantic segmentation tasks.

1. Download the CAMVid dataset.

(a) Dataset Preparation

a.1 Dataset Description

The **CAMVid dataset** consists of urban street scenes, with each pixel labeled according to its corresponding class. The dataset includes training, validation, and test images along with corresponding segmentation masks. The original image resolution is **(960 × 720)**.

a.2 Data Loading and Preprocessing

To ensure proper model input, we apply the following preprocessing steps:

- **Resizing** images to **(480 × 360)**.
- **Normalization** using the mean **[0.485, 0.456, 0.406]** and standard deviation **[0.229, 0.224, 0.225]**.

```
# Define transforms
transform = transforms.Compose([
    transforms.Resize((360, 480)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

- **Encoding segmentation masks**, converting RGB masks into label indices.
- A custom PyTorch **Dataloader** was implemented to handle these transformations efficiently with **batch size: 10**

```
Train Dataset Size: 369
Test Dataset Size: 232
```


```

1  class CamVidDataset(Dataset):
2      def __init__(self, image_dir, mask_dir, class_dict_path, transform=None, preload=True):
3          self.image_dir = image_dir
4          self.mask_dir = mask_dir
5          self.transform = transform
6          self.preload = preload
7          self.class_dict = pd.read_csv(class_dict_path)
8          self.images = os.listdir(image_dir)
9
10         # Load class mapping
11         self.color_to_label = {
12             tuple(self.class_dict.iloc[i, 1:4].astype(int)): i
13             for i in range(len(self.class_dict))
14         }
15
16         self.mask_transform = transforms.Compose([
17             transforms.Resize((360, 480), interpolation=Image.NEAREST)
18         ])
19
20         # Preload data into RAM if preload=True
21         if self.preload:
22             self.preloaded_data = []
23             for img_name in self.images:
24                 self.preloaded_data.append(self.process_image_mask(img_name))
25
26         def __len__(self):
27             return len(self.images)
28
29         def encode_mask(self, mask):
30             mask = np.array(mask, dtype=np.uint8)
31             label_mask = np.zeros(mask.shape[:2], dtype=np.int64)
32             for color, label in self.color_to_label.items():
33                 label_mask[(mask == color).all(axis=-1)] = label
34             return torch.tensor(label_mask, dtype=torch.long)
35
36         def process_image_mask(self, img_name):
37             img_path = os.path.join(self.image_dir, img_name)
38             mask_path = os.path.join(self.mask_dir, img_name.replace('.png', '_L.png'))
39
40             image = Image.open(img_path).convert("RGB")
41             mask = Image.open(mask_path).convert("RGB")
42
43             if self.transform:
44                 image = self.transform(image)
45
46             mask = self.mask_transform(mask)
47             mask = self.encode_mask(mask)
48
49             return image, mask
50
51         def __getitem__(self, idx):
52             return self.preloaded_data[idx] if self.preload else self.process_image_mask(self.images[idx])
53
54         def visualize_sample(self, idx):
55             img_path = os.path.join(self.image_dir, self.images[idx])
56             mask_path = os.path.join(self.mask_dir, self.images[idx].replace('.png', '_L.png'))
57
58             image = Image.open(img_path).convert("RGB")
59             mask = Image.open(mask_path).convert("RGB")
60
61             fig, axes = plt.subplots(1, 2, figsize=(8, 5))
62             axes[0].imshow(image)
63             axes[0].set_title("Image")
64             axes[0].axis("off")
65
66             axes[1].imshow(mask)
67             axes[1].set_title("Mask")
68             axes[1].axis("off")
69
70             plt.show()



```

(b) Class Distribution Visualization

To analyze class frequency in the dataset, we computed the pixel-wise distribution of class labels across the training images.



(c) Image and Mask Visualization



2. SegNet Encoder-Decoder

(a) *Implementation and Training*

a.1 Architecture

The SegNet was implemented by following the architecture specified in the `model_classes.py` file. The encoder consists of five stages, each corresponding to an encoding stage. The layers in the decoder mirror the encoder using transposed convolutions, batch normalization, and activation functions.




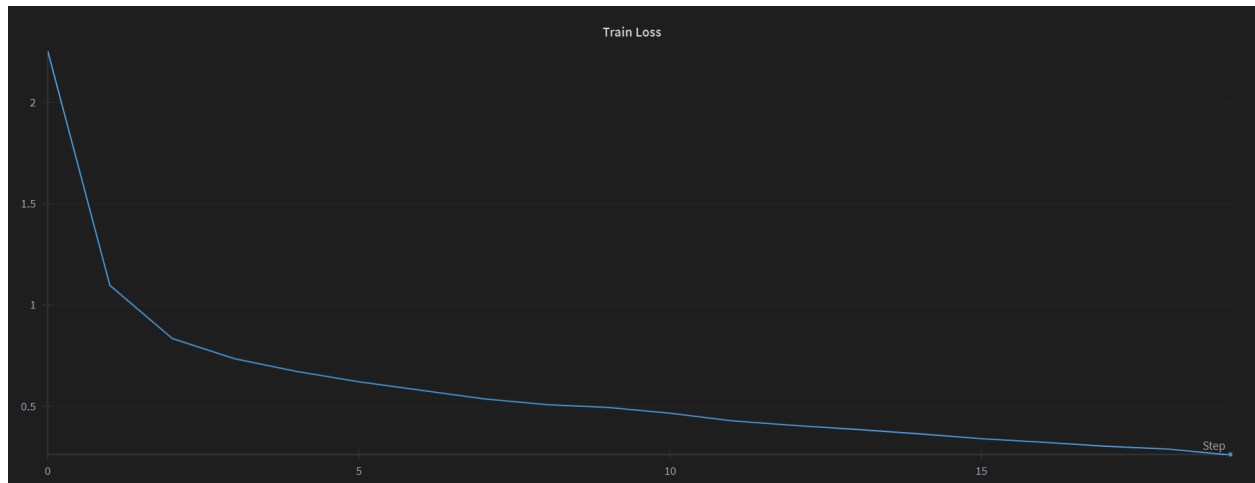
Fig. 2. An illustration of the SegNet architecture. There are no fully connected layers and hence it is only convolutional. A decoder upsamples its input using the transferred pool indices from its encoder to produce a sparse feature map(s). It then performs convolution with a trainable filter bank to densify the feature map. The final decoder output feature maps are fed to a soft-max classifier for pixel-wise classification.

a.2 Training Setup

- **Loss Function:** Cross-entropy loss
- **Optimizer:** Adam optimizer
- **Batch Normalization Momentum:** 0.5
- **Learning rate:** 0.0002
- **Epoch:** 20

segnet	Finished	Add notes	harshu	19m ago	5m 20s	-	20	0.0002	0.26206
--------	----------	-----------	--------	---------	--------	---	----	--------	---------

a.3 Training logs



```
===== TRAINING segnet =====  
Epoch [1/20] -> Train Loss: 2.2558  
Epoch [2/20] -> Train Loss: 1.0993  
Epoch [3/20] -> Train Loss: 0.8367  
Epoch [4/20] -> Train Loss: 0.7373  
Epoch [5/20] -> Train Loss: 0.6747  
Epoch [6/20] -> Train Loss: 0.6237  
Epoch [7/20] -> Train Loss: 0.5817  
Epoch [8/20] -> Train Loss: 0.5400  
Epoch [9/20] -> Train Loss: 0.5114  
Epoch [10/20] -> Train Loss: 0.4970  
Epoch [11/20] -> Train Loss: 0.4687  
Epoch [12/20] -> Train Loss: 0.4307  
Epoch [13/20] -> Train Loss: 0.4084  
Epoch [14/20] -> Train Loss: 0.3887  
Epoch [15/20] -> Train Loss: 0.3666  
Epoch [16/20] -> Train Loss: 0.3428  
Epoch [17/20] -> Train Loss: 0.3251  
Epoch [18/20] -> Train Loss: 0.3056  
Epoch [19/20] -> Train Loss: 0.2917  
Epoch [20/20] -> Train Loss: 0.2621  
===== TRAINING COMPLETED =====
```

a.4 Model Saving

The SegNet decoder was successfully trained using a pre-trained encoder

SegNet_Encoder is saved in **encoder_model.pth**

SegNet_Decoder is saved in **decoder.pth**

(b) Performance Evaluation

Pixel Accuracy: 0.8150

Mean IoU (mIoU): 0.2272

Class-wise Metrics:

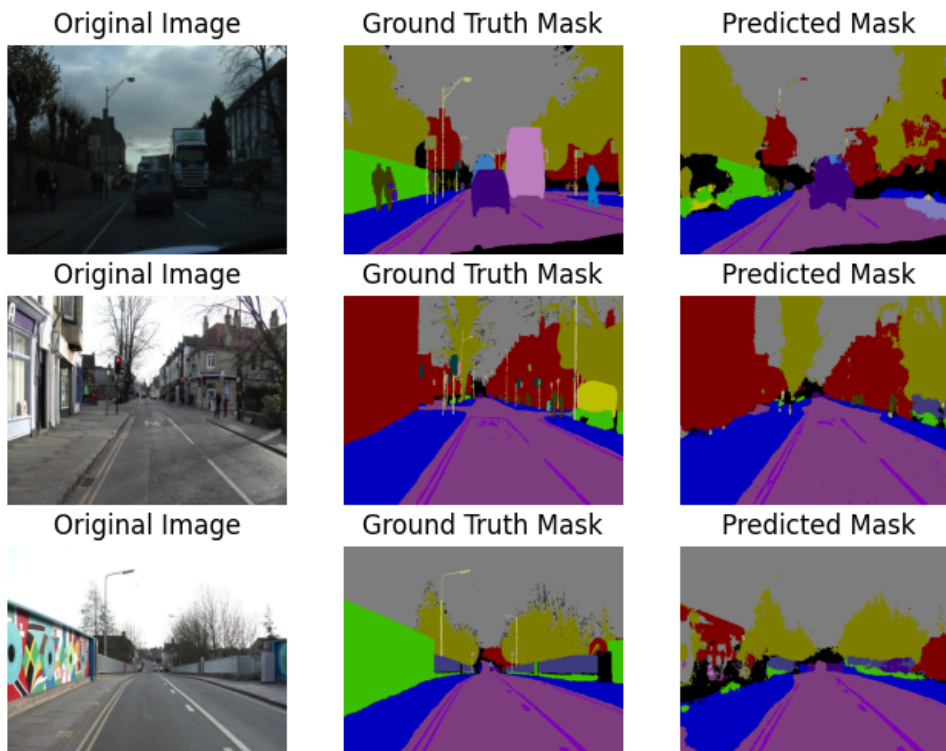
Class 0:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 1:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 2:	IoU=0.3223	Dice=0.4874	Precision=0.7264	Recall=0.3668
Class 3:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 4:	IoU=0.7550	Dice=0.8604	Precision=0.8510	Recall=0.8700
Class 5:	IoU=0.6155	Dice=0.7620	Precision=0.7673	Recall=0.7568
Class 6:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 7:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 8:	IoU=0.0523	Dice=0.0994	Precision=0.4449	Recall=0.0559
Class 9:	IoU=0.2815	Dice=0.4393	Precision=0.4963	Recall=0.3941
Class 10:	IoU=0.4408	Dice=0.6119	Precision=0.7384	Recall=0.5223
Class 11:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 12:	IoU=0.0317	Dice=0.0615	Precision=0.1046	Recall=0.0436
Class 13:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 14:	IoU=0.1922	Dice=0.3225	Precision=0.3902	Recall=0.2748
Class 15:	IoU=0.1480	Dice=0.2579	Precision=0.2204	Recall=0.3106
Class 16:	IoU=0.1188	Dice=0.2124	Precision=0.3126	Recall=0.1609
Class 17:	IoU=0.8758	Dice=0.9338	Precision=0.9376	Recall=0.9300
Class 18:	IoU=0.1430	Dice=0.2502	Precision=0.1621	Recall=0.5479
Class 19:	IoU=0.6839	Dice=0.8123	Precision=0.8031	Recall=0.8217
Class 20:	IoU=0.0052	Dice=0.0104	Precision=0.8651	Recall=0.0052
Class 21:	IoU=0.8945	Dice=0.9443	Precision=0.9310	Recall=0.9580
Class 22:	IoU=0.0585	Dice=0.1105	Precision=0.1739	Recall=0.0809
Class 23:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 24:	IoU=0.1687	Dice=0.2887	Precision=0.6612	Recall=0.1847
Class 25:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 26:	IoU=0.6588	Dice=0.7943	Precision=0.7279	Recall=0.8740
Class 27:	IoU=0.0624	Dice=0.1175	Precision=0.5183	Recall=0.0663
Class 28:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 29:	IoU=0.1131	Dice=0.2033	Precision=0.3519	Recall=0.1429
Class 30:	IoU=0.3189	Dice=0.4836	Precision=0.4197	Recall=0.5704
Class 31:	IoU=0.3309	Dice=0.4972	Precision=0.5536	Recall=0.4513

(c) Visualization and Failure Analysis

c.1 Visualizing $\text{IoU} \leq 0.5$ Cases

For each class, three images where the $0 < \text{IoU} \leq 0.5$ were selected but for some classes there are no examples which might be because class objects are not present in the test dataset. These images show the predicted masks compared to ground truth.

Example Figures:



c.2 Failure Analysis

The model struggles in the following cases:

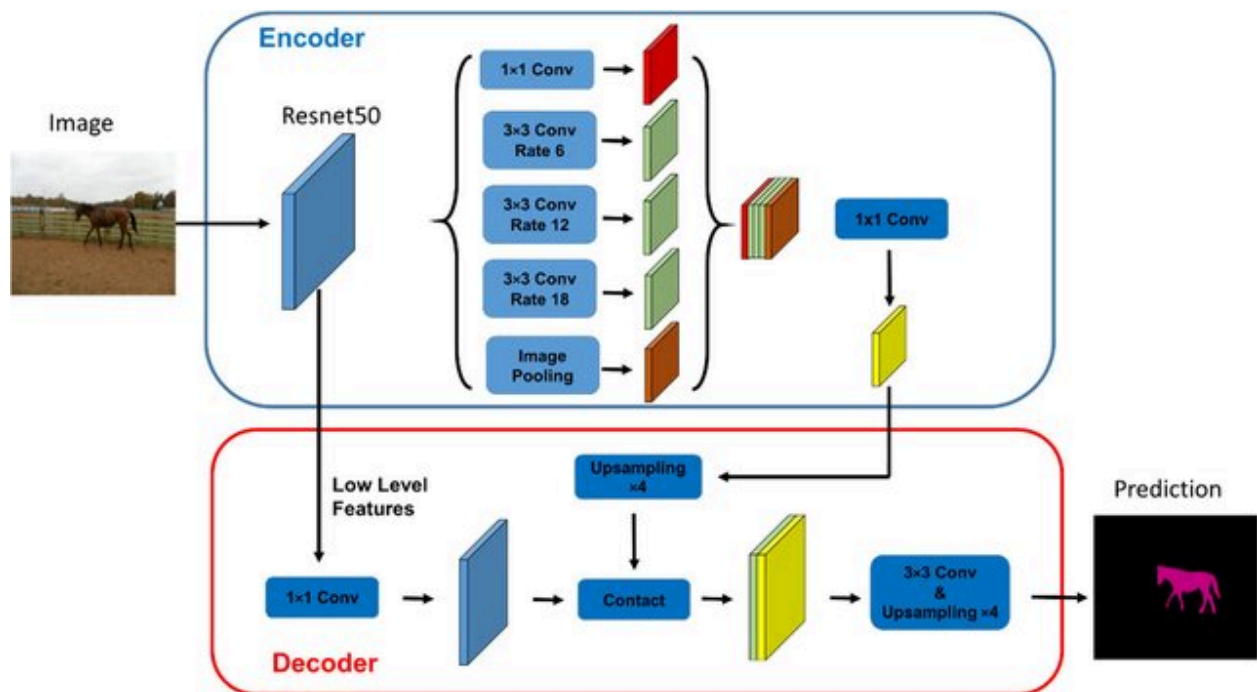
- **Occlusion:** Objects are partially visible.
- **Misclassification:** The model confuses similar-looking objects like bicyclist and pedestrians
- **Environmental Challenges:** Low lighting or complex backgrounds.
- **Small Object Detection:** CamVid contains small, thin objects like poles and signs, which are hard for convolutional layers to segment accurately.
- **Ambiguous Boundaries:** Blurred or unclear boundaries between objects (e.g., road and sidewalk) can confuse the model.
- **SegNet** lacks advanced boundary refinement techniques.

3. DeepLabv3 (ResNet50)

(a) *Implementation and Training*

a.1 Architecture

DeepLabV3, pre-trained on the Pascal VOC dataset, was fine-tuned for CamVid segmentation. The classifier was modified to predict 32 classes.

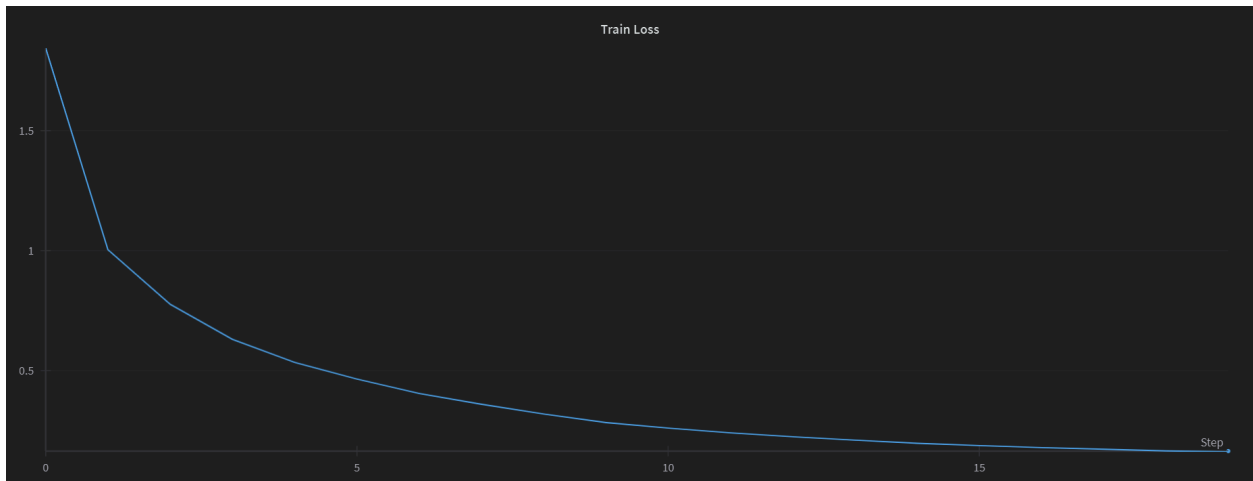


a.2 Training Setup

- **Loss Function:** Cross-entropy loss
- **Optimizer:** Adam optimizer
- **Batch Normalization Momentum:** 0.5
- **Learning rate:** 0.0002
- **Epoch:** 20

•	•	•	deeplabv3	•	Finished	Add notes	harshu	25m ago	14m 24s	-	20	0.0002	0.16492
---	---	---	-----------	---	----------	-----------	--------	---------	---------	---	----	--------	---------

a.3 Training logs



```
===== TRAINING deeplabv3 =====  
Epoch [1/20] -> Train Loss: 1.8444  
Epoch [2/20] -> Train Loss: 1.0057  
Epoch [3/20] -> Train Loss: 0.7784  
Epoch [4/20] -> Train Loss: 0.6327  
Epoch [5/20] -> Train Loss: 0.5366  
Epoch [6/20] -> Train Loss: 0.4672  
Epoch [7/20] -> Train Loss: 0.4074  
Epoch [8/20] -> Train Loss: 0.3625  
Epoch [9/20] -> Train Loss: 0.3218  
Epoch [10/20] -> Train Loss: 0.2863  
Epoch [11/20] -> Train Loss: 0.2632  
Epoch [12/20] -> Train Loss: 0.2437  
Epoch [13/20] -> Train Loss: 0.2273  
Epoch [14/20] -> Train Loss: 0.2134  
Epoch [15/20] -> Train Loss: 0.2003  
Epoch [16/20] -> Train Loss: 0.1907  
Epoch [17/20] -> Train Loss: 0.1823  
Epoch [18/20] -> Train Loss: 0.1756  
Epoch [19/20] -> Train Loss: 0.1684  
Epoch [20/20] -> Train Loss: 0.1649  
===== TRAINING COMPLETED =====
```

a.4 Model Saving

The DeepLabv3 decoder was successfully fine tuned on CamVid

Model's state dict saved as ***deeplabv3.pth***

(b) Performance Evaluation

Pixel Accuracy: 0.8817

Mean IoU (mIoU): 0.3659

Class-wise Metrics:


Class 0:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 1:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 2:	IoU=0.4604	Dice=0.6305	Precision=0.9345	Recall=0.4758
Class 3:	IoU=0.2485	Dice=0.3981	Precision=0.9782	Recall=0.2499
Class 4:	IoU=0.8431	Dice=0.9149	Precision=0.8773	Recall=0.9559
Class 5:	IoU=0.8085	Dice=0.8941	Precision=0.8466	Recall=0.9473
Class 6:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 7:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 8:	IoU=0.1434	Dice=0.2508	Precision=0.5185	Recall=0.1654
Class 9:	IoU=0.5299	Dice=0.6928	Precision=0.6934	Recall=0.6921
Class 10:	IoU=0.3887	Dice=0.5598	Precision=0.7211	Recall=0.4575
Class 11:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 12:	IoU=0.3540	Dice=0.5229	Precision=0.5947	Recall=0.4667
Class 13:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 14:	IoU=0.4276	Dice=0.5990	Precision=0.5526	Recall=0.6540
Class 15:	IoU=0.3900	Dice=0.5611	Precision=0.8549	Recall=0.4176
Class 16:	IoU=0.4044	Dice=0.5760	Precision=0.5887	Recall=0.5637
Class 17:	IoU=0.9145	Dice=0.9553	Precision=0.9463	Recall=0.9646
Class 18:	IoU=0.5162	Dice=0.6809	Precision=0.9262	Recall=0.5384
Class 19:	IoU=0.7920	Dice=0.8839	Precision=0.8349	Recall=0.9390
Class 20:	IoU=0.2989	Dice=0.4602	Precision=0.8748	Recall=0.3123
Class 21:	IoU=0.9137	Dice=0.9549	Precision=0.9401	Recall=0.9702
Class 22:	IoU=0.2391	Dice=0.3860	Precision=0.4827	Recall=0.3215
Class 23:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 24:	IoU=0.5217	Dice=0.6857	Precision=0.7886	Recall=0.6065
Class 25:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 26:	IoU=0.7705	Dice=0.8704	Precision=0.8720	Recall=0.8688
Class 27:	IoU=0.1457	Dice=0.2544	Precision=0.7303	Recall=0.1540
Class 28:	IoU=0.0000	Dice=0.0000	Precision=0.0000	Recall=0.0000
Class 29:	IoU=0.5709	Dice=0.7268	Precision=0.7200	Recall=0.7338
Class 30:	IoU=0.4720	Dice=0.6413	Precision=0.6936	Recall=0.5964
Class 31:	IoU=0.5560	Dice=0.7147	Precision=0.8131	Recall=0.6375

(c) Visualization and Failure Analysis

c.1 Visualizing $\text{IoU} \leq 0.5$ Cases

For each class, three images where the $0 < \text{IoU} \leq 0.5$ were selected but for some classes there are no examples which might be because class objects are not present in the test dataset. These images show the predicted masks compared to ground truth.

Example Figures:



c.2 Failure Analysis

Overall DeepLabv3 is performing better than segnet but it also struggles in some cases:

- **Occlusion:** Objects are partially visible.
- **Misclassification:** The model confuses similar-looking objects like bicyclist and pedestrians
- **Environmental Challenges:** Low lighting or complex backgrounds.
- **DeepLabV3** handles boundaries better but may still struggle with overlapping objects.